

# Stock Price Analysis-Tesla

```
In [695...
#import the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter('ignore')
import yfinance as yf
```

Section 1: The data is non stationary and non mean reverting. The variance is also non stationary and trend is upward and there is a steep change at the end. We have decomposed the data into trend, variance and seasonality and found the similar results

```
In [696...
startdate='2015-01-01'
enddate='2020-12-31'
df=yf.download('TSLA',start=startdate,end=enddate)
```

[\*\*\*\*\*100%\*\*\*\*\*] 1 of 1 completed

```
In [697...
df.head(15)
df.tail(15)
```

```
Out[697...
Open      High      Low      Close      Adj Close      Volume
Date
2020-12-09 653.690002  654.320007  588.000000  604.479980  604.479980  71291200
2020-12-10 574.369995  627.750000  566.340027  627.070007  627.070007  67083200
2020-12-11 615.010010  624.000000  596.799988  609.989990  609.989990  46475000
2020-12-14 619.000000  642.750000  610.200012  639.830017  639.830017  52040600
2020-12-15 643.280029  646.900024  623.799988  633.250000  633.250000  45071500
2020-12-16 628.229980  632.500000  605.000000  622.770020  622.770020  42095800
2020-12-17 628.190002  658.820007  619.500000  655.900024  655.900024  56270100
2020-12-18 668.900024  695.000000  628.539978  695.000000  695.000000  222126200
2020-12-21 666.239990  668.500000  646.070007  649.859985  649.859985  58045300
2020-12-22 648.000000  649.880005  614.229980  640.340027  640.340027  51716000
2020-12-23 632.200012  651.500000  622.570007  645.979980  645.979980  33173000
2020-12-24 642.989990  666.090027  641.000000  661.770020  661.770020  22865600
2020-12-28 674.510010  681.400024  660.799988  663.690002  663.690002  32278600
2020-12-29 661.000000  669.900024  655.000000  665.989990  665.989990  22910800
```

	Open	High	Low	Close	Adj Close	Volume
Date						
2020-12-30	672.000000	696.599976	668.359985	694.780029	694.780029	42846000

In [698... `df=df['Close'].reset_index()`

In [699... `df.head(2)`


Out[699... 

	Date	Close
0	2015-01-02	43.862000
1	2015-01-05	42.018002

In [700... `df['Date'] = pd.to_datetime(df['Date'],infer_datetime_format=True)`  
`df = df.set_index(['Date'])`

## Plot the Trend of Stock Price over the Years

In [701... `plt.title("Closing Prices of Tesla",fontsize=20)`  
`plt.xlabel('Date')`  
`plt.ylabel('Closing price')`  
`plt.plot(df['Close'])`  
`plt.show()`



## Compare with rolling statistics


In [702... `#Mean Stock price on a window of 365 days`  
`df["rolling_avg"] = df["Close"].rolling(window=365).mean()`

```
#Standard Deviation of stock price on 365 days window
df["rolling_std"] = df["Close"].rolling(window=365).std()
```

```
In [703... df=df.reset_index()
```

```
In [704... plt.figure(figsize=(20,10))
plt.title('Trend over the years Compared with yearly moving averages',size=20)
sns.lineplot(data=df, x='Date',y='Close',label='Stock Prize ')
sns.lineplot(data=df,x='Date',y='rolling_avg',label='Yearly Moving Average stock price')
sns.lineplot(data=df,x='Date',y='rolling_std',label='Yearly moving Standard Deviation')
```

```
Out[704... <AxesSubplot:title={'center':'Trend over the years Compared with yearly moving average
s'}, xlabel='Date', ylabel='Close'>
```



## Augmented Dickey–Fuller Test

Checking for non-stationarity

Null Hypothesis: The data is not stationary.

Alternative Hypothesis: The data is stationary.

```
In [705... from statsmodels.tsa.stattools import adfuller,kpss
```

```
In [706... dfctest = adfuller(df['Close'], autolag='AIC')

dfcoutput = pd.Series(dfctest[0:4], index=['Test Statistic','p-value','#Lags Used','Numbe
for key,value in dfctest[4].items():
    dfcoutput['Critical Value (%s)'%key] = value

print(dfcoutput)
```

Test Statistic

6.314977

```

p-value                1.000000
#Lags Used              23.000000
Number of Observations Used  1486.000000
Critical Value (1%)       -3.434758
Critical Value (5%)       -2.863487
Critical Value (10%)      -2.567807
dtype: float64

```

Based on such high P values, we fail to reject the Null so the data is not Stationary

## Kwiatkowski-Phillips-Schmidt-Shin (KPSS)

Here the null hypothesis is that the series is stationary

**Null Hypothesis:** The data is stationary

**Alternate Hypothesis:** The data is not stationary

```
In [707... statistic, p_value, n_lags, critical_values = kpss(df['Close'])
```

```
In [708... print(f'KPSS Statistic: {statistic}')
print(f'p-value: {p_value}')
print(f'num lags: {n_lags}')
print('Critical Values:')
```

```

KPSS Statistic: 2.342020314039496
p-value: 0.01
num lags: 24
Critical Values:

```

Based on a lower p value, we reject the null , the data is non stationary

## Box Cox Transformation

For making variance constant

```
In [709... from scipy import stats
```

```
In [710... Boxcox=list(stats.boxcox(df['Close'])[0])
```

```
In [711... transformed_data, best_lambda = stats.boxcox(df['Close'])
best_lambda
```

```
Out[711... -1.3625118729824754
```

```
In [712... logged=np.log(df['Close'])
```

```
df.insert(len(df.columns), 'logged values',logged)
```


```
In [713... df.insert(len(df.columns), 'Boxcox',transformed_data)
```

```
In [714... df_copy=df.copy()
```

## Plot the values after Box Cox Transformation

```
In [715... plt.figure(figsize=(20,10))
sns.lineplot(data=df,x='Date',y='Boxcox')
```

```
Out[715... <AxesSubplot:xlabel='Date', ylabel='Boxcox'>
```



## Augemented Dicky Fuller Test to check if the data is stationary for box-cox transformed values

**Null Hypothesis:** The data is not stationary.

**Alternative Hypothesis:** The data is stationary.

```
In [716... test= adfuller(df['Boxcox'], autolag='AIC')

output = pd.Series(test[0:4], index=['Test Statistic','p-value','#Lags Used','Number of
for key,value in test[4].items():
    output['Critical Value (%s)'%key] = value

print(output)
```

Test Statistic

-1.285773

```

p-value                0.635656
#Lags Used              0.000000
Number of Observations Used  1509.000000
Critical Value (1%)      -3.434691
Critical Value (5%)      -2.863457
Critical Value (10%)     -2.567791
dtype: float64

```

Based on the high p values, we Fail to reject the Null ,data is not stationary

## Autocorrelation

Since the Mean of data is non Stationary, calculating the first difference of data, the 3rd difference and the 6th difference

Calculating values of AutoCorrelation, a measure of how correlated time series data is at a given point in time with past values

### First lag autocorrelation

```

In [717... autocorrelation_lag1=df['Close'].autocorr(lag=30) ## Lag taken for 30 days
print("1 month Lag:", autocorrelation_lag1)

```

1 month Lag: 0.9624867552011044

### Second lag autocorrelation value of Data: 0.9601946480498523

```

In [718... autocorrelation_lag2=df['Close'].autocorr(lag=60) ## Lag taken for 60 days
print("2 month Lag:", autocorrelation_lag2)

```

2 month Lag: 0.9456946549894001

### Third lag autocorrelation value of Data: 0.8956753113926396

```

In [719... autocorrelation_lag3=df['Close'].autocorr(lag=120) ## Lag taken for 120 days
print(" 3 month lag:", autocorrelation_lag3)

```

3 month lag: 0.8959107176901502

Inference : even after giving lag of 3 months the data is highly correlated


## ACF/PACF Plots for closing price value

```

In [720... from statsmodels.graphics.tsaplots import plot_acf
fig1=plt.figure(figsize=(30,10))
fig1=plot_acf(df['Close'])

```


<Figure size 2160x720 with 0 Axes>



In [721]...

```
from statsmodels.graphics.tsaplots import plot_pacf
fig=plt.figure(figsize=(30,10))
fig=plot_pacf(df['Close'],lags=30)
```

&lt;Figure size 2160x720 with 0 Axes&gt;




From ACF and PACF plots it looks like AR of 2nd order

## ACF PACF plot using 1st Lag-(1 month lag approx 30 days)

In [722]...

```
df['lagprice'] = df['Close'].shift(30)
df['1st Differencing']=df['lagprice']-df['Close']
df.dropna(inplace=True)
fig4=plt.figure(figsize=(30,10))
fig4=plot_acf(df['1st Differencing'])
```


&lt;Figure size 2160x720 with 0 Axes&gt;



In [723...

```
fig=plt.figure(figsize=(30,10))
fig=plot_pacf(df['1st Differencing'],lags=30)
```

&lt;Figure size 2160x720 with 0 Axes&gt;




## ACF PACF plot using 2nd Lag

In [724...

```
df['lagprice'] = df['Close'].shift(60)
df['1st Differencing']=df['lagprice']-df['Close']
df.dropna(inplace=True)
fig4=plt.figure(figsize=(30,10))
fig4=plot_acf(df['1st Differencing'])
```

&lt;Figure size 2160x720 with 0 Axes&gt;






In [725...

```
fig=plt.figure(figsize=(30,10))
fig=plot_pacf(df['1st Differencing'],lags=60)
```

&lt;Figure size 2160x720 with 0 Axes&gt;




## ACF PACF plot using 3rd Lag- approx 3 months

In [726...

```
df['lagprice'] = df['Close'].shift(90)
df['1st Differencing']=df['lagprice']-df['Close']
df.dropna(inplace=True)
fig4=plt.figure(figsize=(30,10))
fig4=plot_acf(df['1st Differencing'])
```


&lt;Figure size 2160x720 with 0 Axes&gt;



In [727...

```
fig=plt.figure(figsize=(30,10))
fig=plot_pacf(df['1st Differencing'],lags=90)
```

&lt;Figure size 2160x720 with 0 Axes&gt;



In [728...

```
dfctest1 = adfuller(df['1st Differencing'], autolag='AIC')

dfoutput = pd.Series(dfctest1[0:4], index=['Test Statistic','p-value','#Lags Used','Numb
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%s)'%key] = value

print(dfoutput)
```


Test Statistic	0.235194
p-value	0.974168
#Lags Used	19.000000
Number of Observations Used	976.000000
Critical Value (1%)	-3.434758
Critical Value (5%)	-2.863487
Critical Value (10%)	-2.567807
dtype:	float64

Inference: Even after taking a lag of 3 months , the p value is high and is greater than 0.05 so the data is not stationary

## ACF PACF plots for Lagged differencing on BoxCox transformed Values


```
In [729...
df['lagprice'] = df['Boxcox'].shift(3)
df['3rd lag Differencing']=df['lagprice']-df['Close']
df.dropna(inplace=True)
fig4=plt.figure(figsize=(30,10))
fig4=plot_acf(df['3rd lag Differencing'])
```

<Figure size 2160x720 with 0 Axes>



```
In [730...
fig5=plt.figure(figsize=(30,10))
fig5=plot_pacf(df['3rd lag Differencing'],lags=30)
```

<Figure size 2160x720 with 0 Axes>



from the ACF and PACF plots it looks like AR of 2nd order

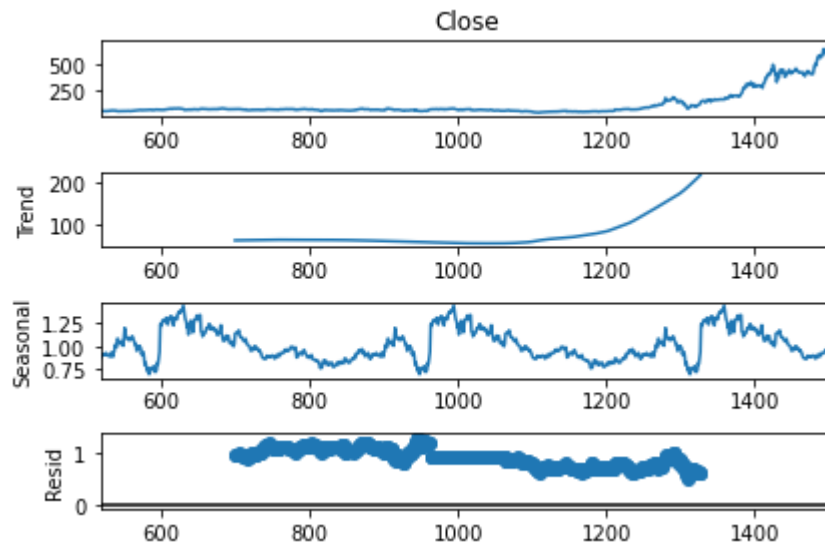
# Trend Decomposition

```
In [731... from statsmodels.tsa.seasonal import seasonal_decompose
```

```
In [732... decompose=seasonal_decompose(df['Close'],model='multiplicative',period=365) ## period i
```

```
In [733... fig3=plt.figure(figsize=(30,20))
fig3=decompose.plot()
```

<Figure size 2160x1440 with 0 Axes>



Inference: Yes, there are seasonal fluctuations in the data as can be seen from the seasonal spikes in the figure above but overall there is no seasonality since there is a high spike at the end

## Section 2

### Fitting Several Arima Models and the using Auto Arima to obtain the best model

```
In [734... import pmdarima
from pmdarima.arima import auto_arima
```

```
In [735... from statsmodels.tsa.arima.model import ARIMA
import statsmodels.api as sm
```

```
In [736... model1=ARIMA(df_copy['Boxcox'],order=(2,0,1))
model1 = model1.fit()
model1.summary()
print(model1.bic)
```

-15497.180506688877

```
In [737... model2=ARIMA(df_copy['Boxcox'],order=(1,1,1))
model2 = model2.fit()
model2.summary()
print(model2.bic)
```

-22753.742733672225

```
In [738... model3=ARIMA(df_copy['Boxcox'],order=(2,1,3))
model3 = model3.fit()
model3.summary()
print(model3.bic)
```


-22731.962556627073

Based on the above findings, we select the ARIMA(1,1,1) model as it has the least AIC and BIC

```
In [739... model_best=ARIMA(df_copy['Boxcox'],order=(1,1,1))
```

```
In [740... model_fit_best = model_best.fit()
residuals_best = pd.DataFrame(model_fit_best.resid)
residuals_best.plot()
```

Out[740... <AxesSubplot:>



```
In [741... residuals_best.describe()
```

```
Out[741... 0
count  1510.000000
mean    0.000486
std     0.018778
min    -0.000858
```

	0
25%	-0.000045
50%	0.000002
75%	0.000060
max	0.729689

## Auto Arima to find the best fit model

In [742...

```
import pmdarima as pm
```

In [743...

```
model_autoarima = pm.auto_arima(df_copy['Boxcox'], start_p=0, start_q=0,
                                test='adf',          # use adftest to find optimal 'd'
                                max_p=5, max_q=5,    # maximum p and q
                                m=5,                # frequency of series
                                d=None,              # let model determine 'd'

                                start_P=0,
                                D=0,
                                trace=True,
                                error_action='ignore',
                                suppress_warnings=True,
                                stepwise=True)
```

Performing stepwise search to minimize aic

```
ARIMA(0,1,0)(0,0,1)[5] intercept : AIC=-22769.073, Time=1.09 sec
ARIMA(0,1,0)(0,0,0)[5] intercept : AIC=-22770.545, Time=0.74 sec
ARIMA(1,1,0)(1,0,0)[5] intercept : AIC=-22767.703, Time=1.06 sec
ARIMA(0,1,1)(0,0,1)[5] intercept : AIC=-22767.648, Time=1.63 sec
ARIMA(0,1,0)(0,0,0)[5]          : AIC=-22771.844, Time=0.11 sec
ARIMA(0,1,0)(1,0,0)[5] intercept : AIC=-22769.058, Time=0.87 sec
ARIMA(0,1,0)(1,0,1)[5] intercept : AIC=-22767.284, Time=1.63 sec
ARIMA(1,1,0)(0,0,0)[5] intercept : AIC=-22769.267, Time=0.53 sec
ARIMA(0,1,1)(0,0,0)[5] intercept : AIC=-22769.208, Time=1.19 sec
ARIMA(1,1,1)(0,0,0)[5] intercept : AIC=-22768.237, Time=2.54 sec
```

Best model: ARIMA(0,1,0)(0,0,0)[5]

Total fit time: 11.412 seconds

In [744...

```
model_autoarima.summary()
```

Out[744...

### SARIMAX Results

<b>Dep. Variable:</b>	y	<b>No. Observations:</b>	1510
<b>Model:</b>	SARIMAX(0, 1, 0)	<b>Log Likelihood</b>	11386.922
<b>Date:</b>	Sat, 05 Feb 2022	<b>AIC</b>	-22771.844
<b>Time:</b>	17:26:06	<b>BIC</b>	-22766.525
<b>Sample:</b>	0	<b>HQIC</b>	-22769.863
	- 1510		

Covariance Type:

opg

	coef	std err	z	P> z	[0.025	0.975]
<b>sigma2</b>	1.633e-08	2.89e-10	56.537	0.000	1.58e-08	1.69e-08

**Ljung-Box (L1) (Q):** 0.72    **Jarque-Bera (JB):** 2738.83

**Prob(Q):** 0.40    **Prob(JB):** 0.00

**Heteroskedasticity (H):** 0.56    **Skew:** -0.28

**Prob(H) (two-sided):** 0.00    **Kurtosis:** 9.58


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Ljung-Box test for residuals

In [745...

```
# line plot of residuals
model=ARIMA(df['Boxcox'],order=(0,1,0))
residuals = pd.DataFrame(model_fit_best.resid)
residuals.plot()
plt.show()
```



**H0:** The residuals are independently distributed.

**HA:** The residuals are not independently distributed; they exhibit serial correlation.

In [746...

```
import statsmodels.api as sm
sm.stats.acorr_ljungbox(residuals, lags=[5], return_df=True)
```


Out[746...

	lb_stat	lb_pvalue
<b>5</b>	0.000238	1.0

**Inference : We cannot reject null hypothesis which means residuals are independently distributed**

In [747...

```
# line plot of residuals
model=ARIMA(df_copy['Boxcox'],order=(0,1,0))
residuals = pd.DataFrame(model_fit_best.resid)
residuals.plot()
plt.show()
```



In [761...

```
from scipy.special import boxcox, inv_boxcox
prediction1=model_fit_best.predict(start=0, end=1509)
pred1=inv_boxcox(prediction1,-1.36)
```

In [762...

```
prediction=model_fit_best.predict(start=0, end=1509)
from sklearn.metrics import mean_squared_error
np.sqrt(mean_squared_error(df_copy['Close'],pred1))
```

Out[762...


106.5792299366475

In [763...

```
from statsmodels.graphics.tsaplots import plot_pacf
fig=plt.figure(figsize=(30,10))
fig=plot_pacf(df['Close'],lags=30)
```

<Figure size 2160x720 with 0 Axes>






## ACF PACF plots for residual

In [764...

```
fig4=plt.figure(figsize=(30,10))
fig4=plot_acf(residuals)
```


<Figure size 2160x720 with 0 Axes>



In [765...

```
fig5=plt.figure(figsize=(30,10))
fig5=plot_pacf(residuals,lags=30)
```

<Figure size 2160x720 with 0 Axes>



From ACF PACF figures, it can be inferred that there is no autocorrelation in the residuals and it is a white noise process

RMSE value is 106.58

Predicted values for the next 5 periods

```
In [766... prediction2=model_fit_best.predict(start=1510, end=1514)
pred2=inv_boxcox(prediction2, -1.362)
```


```
In [767... pred2
```

```
Out[767... 1510    261.541310
1511    261.583308
1512    261.606839
1513    261.620021
1514    261.627405
Name: predicted_mean, dtype: float64
```

The Predicted values are reasonable and comparable to the original values

```
In [768... plt.figure(figsize=(20,10))
plt.plot(df['Close'],label='actual')
plt.plot(pred1,label='predicted')
plt.legend()
```

```
Out[768... <matplotlib.legend.Legend at 0x25148adb100>
```



**Yes in sample predicted values follow approximately the same trend as the original data**

In [ ]:

In [ ]: